

```
In [25]: import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: data = pd.read_csv('uber.csv')
```

```
In [3]: df = data.copy()
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.73
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.72
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.74
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.79
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.74

```
In [9]: print (df.shape)
print (df.columns)

(200000, 7)
Index(['fare_amount', 'pickup_datetime', 'pickup_longitude', 'pickup_latitude',
      'dropoff_longitude', 'dropoff_latitude', 'passenger_count'],
      dtype='object')
```

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount           200000 non-null  float64
1   pickup_datetime       200000 non-null  object
2   pickup_longitude      200000 non-null  float64
3   pickup_latitude       200000 non-null  float64
4   dropoff_longitude     199999 non-null  float64
5   dropoff_latitude      199999 non-null  float64
6   passenger_count       200000 non-null  int64
dtypes: float64(5), int64(1), object(1)
memory usage: 10.7+ MB
```

```
In [12]: df["pickup_datetime"]=pd.to_datetime(df['pickup_datetime'])
```

In [13]: `df.head()`

Out[13]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.738354
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.728225
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.740770
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.790844
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.744085

In [14]: `df.describe()`

Out[14]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000
mean	11.359955	-72.527638	39.935885	-72.525292	39.923890
std	9.901776	11.437787	7.720539	13.117408	6.794829
min	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.985513
25%	6.000000	-73.992065	40.734796	-73.991407	40.733823
50%	8.500000	-73.981823	40.752592	-73.980093	40.753042
75%	12.500000	-73.967154	40.767158	-73.963658	40.768001
max	499.000000	57.418457	1644.421482	1153.572603	872.697628

In [15]: `df.isnull().sum()`

Out[15]:

```
fare_amount      0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude 1
dropoff_latitude  1
passenger_count  0
dtype: int64
```


In [17]: `df.dropna(inplace=True)`
`print(df.isnull().sum())`

```
fare_amount      0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude 0
dropoff_latitude  0
passenger_count  0
dtype: int64
```

```
In [18]: import matplotlib.pyplot as plt  
import seaborn as sns  
%matplotlib inline
```

```
In [26]: sns.distplot(df['fare_amount'])
```


```
Out[26]: <AxesSubplot:xlabel='fare_amount', ylabel='Density'>
```



Above we can see some fares having negative values


```
In [27]: sns.distplot(df['pickup_latitude'])
```

```
Out[27]: <AxesSubplot:xlabel='pickup_latitude', ylabel='Density'>
```



```
In [28]: sns.distplot(df['pickup_longitude'])
```


```
Out[28]: <AxesSubplot:xlabel='pickup_longitude', ylabel='Density'>
```



Negative and Positive values are exceeding far behind the real limit.


```
In [29]: sns.distplot(df['dropoff_longitude'])
```

```
Out[29]: <AxesSubplot:xlabel='dropoff_longitude', ylabel='Density'>
```



```
In [30]: sns.distplot(df['dropoff_latitude'])
```

```
Out[30]: <AxesSubplot:xlabel='dropoff_latitude', ylabel='Density'>
```



```
In [32]: import calendar
df['day']=df['pickup_datetime'].apply(lambda x:x.day)
df['hour']=df['pickup_datetime'].apply(lambda x:x.hour)
df['weekday']=df['pickup_datetime'].apply(lambda x:calendar.day_name[x.weekday()])
df['month']=df['pickup_datetime'].apply(lambda x:x.month)
df['year']=df['pickup_datetime'].apply(lambda x:x.year)
```

```
In [33]: df.head()
```

```
Out[33]:
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.738354
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.728225
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.740770
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.790844
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.744085

```
In [34]: df.weekday = df.weekday.map({'Sunday':0, 'Monday':1, 'Tuesday':2, 'Wednesday':3, 'Thursday':4, 'Friday':5, 'Saturday':6})
```

In [35]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 199999 entries, 0 to 199999
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount           199999 non-null  float64
1   pickup_datetime       199999 non-null  datetime64[ns, UTC]
2   pickup_longitude      199999 non-null  float64
3   pickup_latitude       199999 non-null  float64
4   dropoff_longitude     199999 non-null  float64
5   dropoff_latitude      199999 non-null  float64
6   passenger_count       199999 non-null  int64
7   day                   199999 non-null  int64
8   hour                  199999 non-null  int64
9   weekday               199999 non-null  int64
10  month                 199999 non-null  int64
11  year                  199999 non-null  int64
dtypes: datetime64[ns, UTC](1), float64(5), int64(6)
memory usage: 19.8 MB
```

In [37]: df=df[df['passenger_count']<=8]

In [38]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 199998 entries, 0 to 199999
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount           199998 non-null  float64
1   pickup_datetime       199998 non-null  datetime64[ns, UTC]
2   pickup_longitude      199998 non-null  float64
3   pickup_latitude       199998 non-null  float64
4   dropoff_longitude     199998 non-null  float64
5   dropoff_latitude      199998 non-null  float64
6   passenger_count       199998 non-null  int64
7   day                   199998 non-null  int64
8   hour                  199998 non-null  int64
9   weekday               199998 non-null  int64
10  month                 199998 non-null  int64
11  year                  199998 non-null  int64
dtypes: datetime64[ns, UTC](1), float64(5), int64(6)
memory usage: 19.8 MB
```

In [40]: from sklearn.model_selection import train_test_split

In [41]: x=df.drop("fare_amount", axis=1)

In [42]: y=df['fare_amount']

In [43]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,ra

In [44]: `x_train.head()`

Out[44]:

	<code>pickup_datetime</code>	<code>pickup_longitude</code>	<code>pickup_latitude</code>	<code>dropoff_longitude</code>	<code>dropoff_latitude</code>
80768	2009-02-22 01:12:00+00:00	-73.983703	40.725752	-73.972000	40.793881
111783	2009-03-07 14:49:00+00:00	-73.961175	40.760667	-73.976507	40.747570
24615	2011-03-17 11:51:08+00:00	-73.947784	40.783111	-73.955408	40.779401
46932	2010-01-15 07:01:38+00:00	-73.980596	40.733797	-73.972092	40.747291
86655	2014-06-28 19:25:00+00:00	-73.963035	40.758380	-73.987877	40.745471

In [45]: `x_test.head()`

Out[45]:

	<code>pickup_datetime</code>	<code>pickup_longitude</code>	<code>pickup_latitude</code>	<code>dropoff_longitude</code>	<code>dropoff_latitude</code>
13588	2013-06-25 22:32:00+00:00	-73.982810	40.771687	-73.977065	40.763201
29803	2011-02-20 20:16:00+00:00	-73.991985	40.725763	-73.995762	40.759791
138266	2011-08-09 20:15:00+00:00	-73.989458	40.741665	-73.983463	40.758841
82856	2011-11-17 18:49:03+00:00	-73.973200	40.748100	-73.973500	40.748201
162748	2012-08-26 04:37:52+00:00	-73.856011	40.824512	-73.981732	40.761751

In [50]: `x_train['pickup_datetime'] = pd.to_numeric(pd.to_datetime(x_train['pickup_datetime']))`
`x_test['pickup_datetime'] = pd.to_numeric(pd.to_datetime(x_test['pickup_datetime']))`

In [51]: `x_train.shape`

Out[51]: (159998, 11)

In [52]: `x_test.shape`

Out[52]: (40000, 11)

In [53]: `from sklearn.linear_model import LinearRegression`

In [54]: `lrmodel=LinearRegression()`
`lrmodel.fit(x_train, y_train)`

Out[54]: LinearRegression()

In [55]: `predictedvalues = lrmodel.predict(x_test)`

```
In [56]: from sklearn.metrics import mean_squared_error  
lrmse = np.sqrt(mean_squared_error(predictedvalues, y_test))  
print("RMSE value for Linear regression is", lrmse)
```

RMSE value for Linear regression is 9.922284743566864

```
In [57]: from sklearn.ensemble import RandomForestRegressor  
rfrmodel = RandomForestRegressor(n_estimators=100, random_state=101)
```

```
In [58]: rfrmodel.fit(x_train,y_train)  
rfrmodel_pred= rfrmodel.predict(x_test)
```

```
In [59]: rfrmodel_rmse=np.sqrt(mean_squared_error(rfrmodel_pred, y_test))  
print("RMSE value for Random forest regression is ",rfrmodel_rmse)
```

RMSE value for Random forest regression is 4.79773791418236

```
In [ ]:
```